

# Accelerating DI Ce via FPGAs

## Honeywell FM&T Sponsored Research

Presenter: Keaten N. Stokke

Advisor: Dr. David Andrews

September 6th, 2019

# Thank You Honeywell!

- Provided funding for 1 Ph.D. and 1 Master's student for 3 semesters
- Learned a tremendous amount of engineering skills
- Submitted our first ever paper to the ReConFig conference
- Granted me the ability and inspiration to pursue graduate school

**Special shoutout to Mr. Dennis Stanley!!!**

# Overview

- **Objective**
- Results
- Challenges
- Additional & Future Works
- Conclusion
- Project Decomposition
  - Phase 1: Understanding DICE
  - Phase 2: Identifying Key Functions
  - Phase 3: Architecture Design
  - Phase 4: Client [PC]
  - Phase 5: Server [FPGA - PS]
  - Phase 6: Application [FPGA - PL]
  - Phase 7: Testing & Improvements

# Objective

- Investigate hardware acceleration of the Digital Image Correlation engine (DICE) software from Sandia National Laboratories
- The use of FPGAs gave us a few possible approaches to this problem
  - 1) Vivado HLS - “C to Gates”
  - 2) Convert DICE to Verilog
  - 3) Port to the HMC

**We chose option two and opted to convert key functions in DICE to Verilog!**

# Objective


- We took the route of hardcoding key DICE functions from C++ to Verilog because it gave us more options
  - Allowed us to target FPGAs
  - Allowed for the potential use of the HMC by Micron
- The Hybrid Memory Cube (HMC) would be a viable option for the future to eliminate Ethernet data transfer
- DICE source code required our full attention

# Overview

- Objective
- **Results**
- Challenges
- Additional & Future Works
- Conclusion
- Project Decomposition
  - Phase 1: Understanding DICE
  - Phase 2: Identifying Key Functions
  - Phase 3: Architecture Design
  - Phase 4: Client [PC]
  - Phase 5: Server [FPGA - PS]
  - Phase 6: Application [FPGA - PL]
  - Phase 7: Testing & Improvements

# Results

- Created a working end-to-end DICE hardware accelerator prototype
  - A Python script is run on the PC to transmit available frames to FPGA
  - The FPGAs LWIP server receives the frame, converts to IEEE 754
  - The processed frame is then written into BRAM, a start signal is sent
  - Correlation on the PL starts, server requests new frame if needed
  - Results are saved to BRAM, read from the server, and sent to the client
  - Client receives results, converts and formats, writes to result .txt files

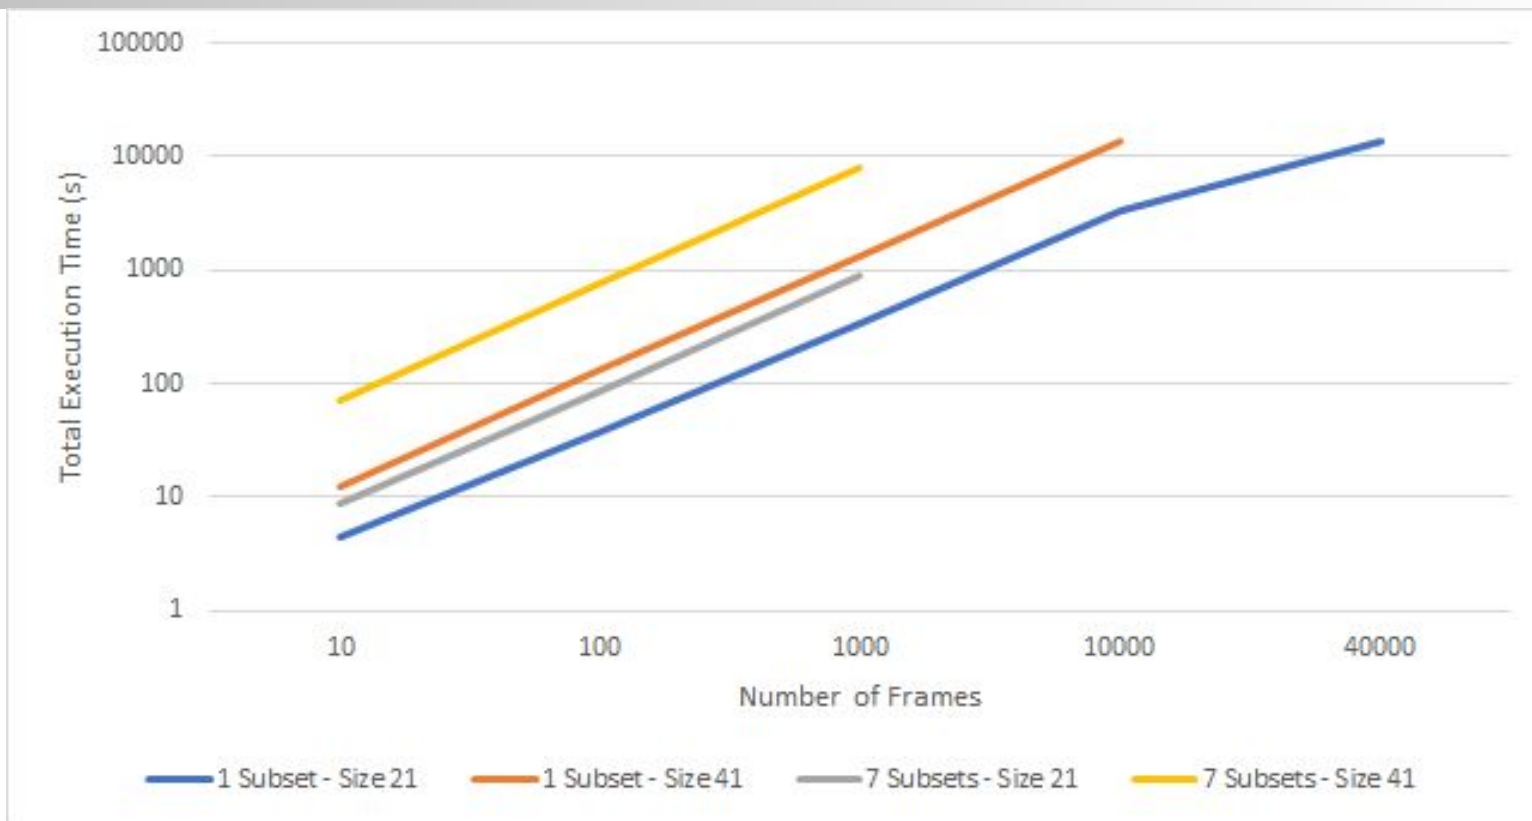


# Results

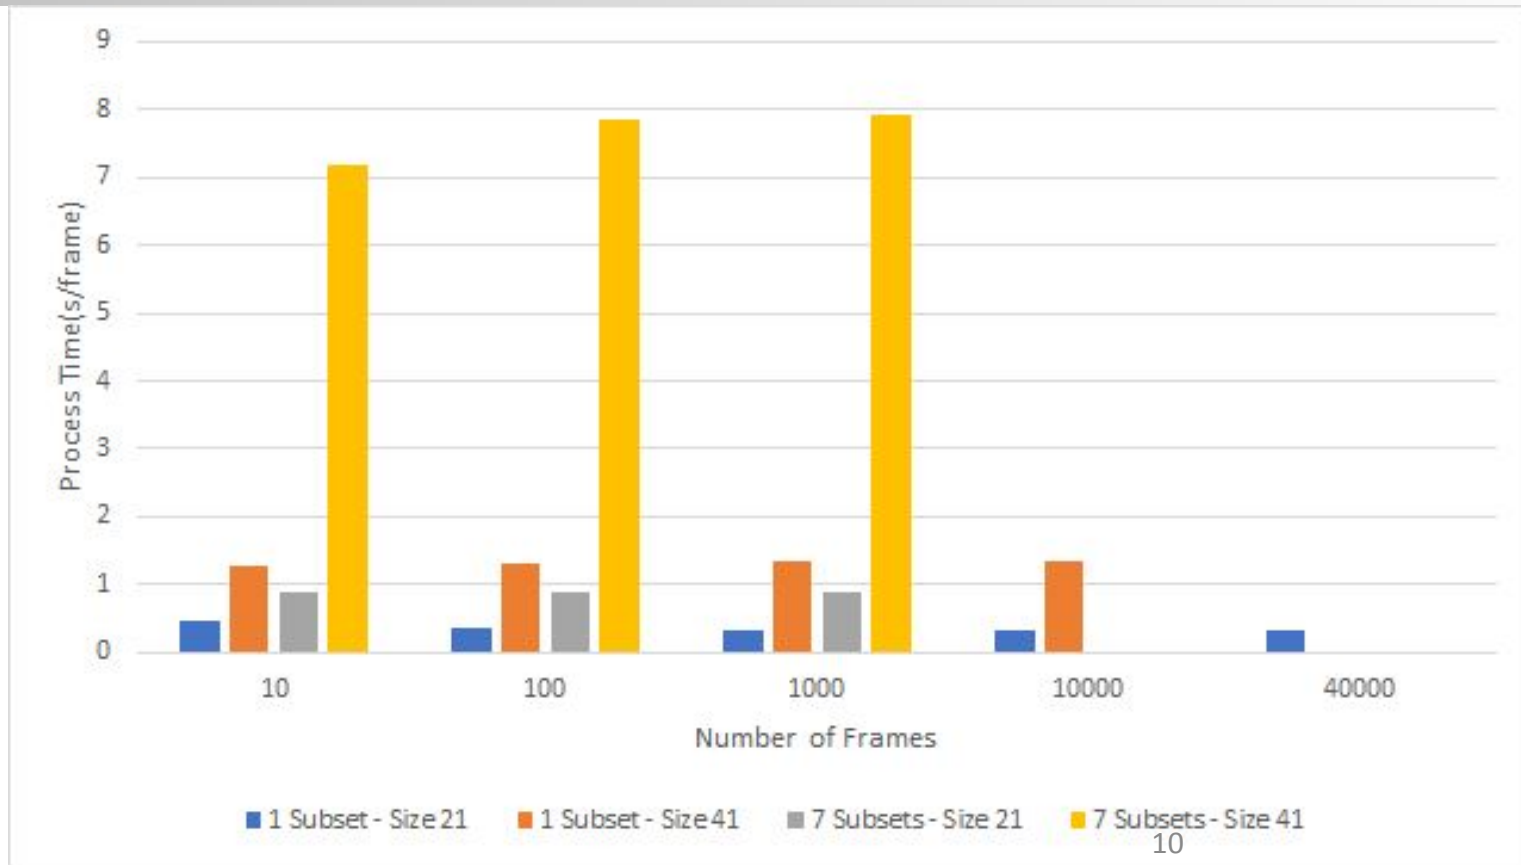
- Ran successful tests of this system with real images
  - Test with 1 subset
    - 40,000 frames, 21x21 subset (control)
    - 10,000 frames, 41x41 subset (quadrupled)
  - Test with 7 subsets
    - 1,000 frames, 21x21 subset (control)
    - 1,000 frames, 41x41 subset (quadrupled)



# Results




# Results



# Results

Resource	Utilization	Available	Utilization %
LUT	37499	230400	16.28
LUTRAM	2787	101760	2.74
FF	33362	460800	7.24
BRAM	279	312	89.42
URAM	96	96	100.00
DSP	11	1728	0.64
IO	3	360	0.83
BUFG	4	544	0.74
MMCM	1	8	12.50

## ZCU104 Resource Utilization



# Results

Test Case: 2 Frames (448x232), 1 Subset (21x21)					
Test Run	DICe GUI	Simulation (Current Design)	Simulation (Pipelined Design)	Demo (Total)	Demo (Ethernet)
Execution Time:	116 ms	138 ms	55 ms	947 ms	933 ms

- Our current demo is 8.16x slower than the DICe GUI
  - Correlation: 14%, Client: 1.5%, **Ethernet: 84.5%**
  - Our simulations are only 1.19x slower than the DICe GUI
- With a pipelined design, it could be 2.1x faster than the DICe GUI
  - We would need more BRAM resources to save the gradients

# Results

## Ethernet Capability

FPGA	Processor	Ethernet IP	Speed (Mbps)	Actual Speed (Mbps)	Tested Speed (Mbps)
Virtex-7 [VC707]	MicroBlaze	AXI Ethernet Subsystem (SGMII)	1000	X	X
Kintex-7 [KC705]	MicroBlaze	AXI Ethernet Lite (MII)	1000	56	8
UltraScale+ MPSoC [ZCU104]	Zynq	PHY (RGMII)	1000	950	+/- 50 (*500)

- After many attempts, Virtex-7 failed to work and Kintex-7 was slow  
**Zynq is the way to go!**

# Results

## KC705 vs. ZCU104

Total Simulated Execution Time vs. DICE GUI

- ZCU104: 12x faster than KC705
  - Ethernet speed can theoretically be 19x faster [ $\sim 50\text{Mbps} \rightarrow \sim 950\text{Mbps}$ ]
- Note that the simulated FPGA design is 6.4x faster than the GUI (when correlating 2 frames)

Test Case: 64x48 Design	Total Transfer Time	Total FPGA Correlation Time	Total GUI Correlation Time	Total GUI Execution Time	FPGA vs. GUI
KC705	1.25 s	0.0025 s	0.016 s	0.058 s	25x slower
ZCU104	0.104 s	0.0025 s	0.016 s	0.058 s	5.6x slower

# Results

## Verification

- Two major drawbacks exist with the current state of our application
  - 1) Results are unverified
  - 2) Some of the results haven't been implemented (beta, sigma)
- The results we do have are not correct when compared to the outputs of the DICE GUI, we have been working to correct this
- Currently, we produce the results Displacement X, Displacement Y, and Rotation Z.


# Results

- Created a library of floating-point operations that resulted in a paper
  - Submitted to the ReConFig conference
  - Functions: addition, subtraction, multiplication, division, sine, cosine, **arcsine**, **arccosine**
- Aimed at reducing latency for each operation
  - Previous works emphasize on pipelining which wasn't useful
- Novel work is based on Finite-State Machines (FSMs)
  - Operations can be locally embedded in custom IP
  - Runs at 150 MHz, improved functions work at 200 MHz




# Results

## Floating-Point Library: Delay Comparison




# Results

## Floating-Point Library: LUTs Comparison



# Results

## Floating-Point Library: FFs Comparison



# Results

## Floating-Point Arithmetic Functions

Function	# Clks	Freq (MHz)	Delay (ns)	DSP	# Slices	Power (W)	(F) / (D*S*M)
Adder	4	150	23.1	-	40	0.37	129.8
Subtractor	5	150	29.7	-	44	0.37	84.1
Multiplier	5	150	29.7	✓	56	0.37	72.1
Divider	30	150	194.7	-	39	0.37	15.4

# Results

## Floating-Point Trigonometric Functions

Function	# Clks	Freq (MHz)	Delay (ns)	DSP	# Slices	Power (W)	(F) / (D*S*M)
<b>Sine</b>	<b>92</b>	<b>150</b>	<b>603.9</b>	<b>✓</b>	<b>351</b>	<b>0.39</b>	<b>0.53</b>
<b>Cosine</b>	<b>86</b>	<b>150</b>	<b>564.3</b>	<b>✓</b>	<b>328</b>	<b>0.39</b>	<b>0.61</b>
<b>Arcsine</b>	<b>97</b>	<b>150</b>	<b>639.6</b>	<b>✓</b>	<b>365</b>	<b>0.39</b>	<b>0.48</b>
<b>Arccosine</b>	<b>105</b>	<b>150</b>	<b>689.7</b>	<b>✓</b>	<b>339</b>	<b>0.39</b>	<b>0.48</b>

# Results

## What comes in the box...

- Documentation for all IP
  - Each IP comes with its own documentation that provides a description of what its function is
- FPGA security report
  - Additionally, a report is provided on FPGA security that highlights the capabilities of the ZCU104
- Everything is available on GitHub
  - [github.com/keatenstokke/DICE](https://github.com/keatenstokke/DICE)

# Results

## Knowledge Gained

- How to use the Vivado tools (2015.4 and 2018.3)
- Working with series-7 FPGAs (Kintex-7 and Virtex-7)
- Working with Zynq FPGAs (Zybo and UltraScale+ MPSoC)
  - How to target multiple cores simultaneously for processing
  - How to transmit data from PS → PL
- How to generate processor interrupts
- How to create and design custom IPs
- How to interface with memories (BRAM, DRAM, FLASH)
- Learned to setup Ethernet data transmission
- Learned how to interface with the LWIP
- Debugging HDL with VIO

**... And the list goes on!**

# Overview

- Objective
- Results
- **Challenges**
- Additional & Future Works
- Conclusion
- Project Decomposition
  - Phase 1: Understanding DICE
  - Phase 2: Identifying Key Functions
  - Phase 3: Architecture Design
  - Phase 4: Client [PC]
  - Phase 5: Server [FPGA - PS]
  - Phase 6: Application [FPGA - PL]
  - Phase 7: Testing & Improvements



# Challenges

- Assembling the design's infrastructure
  - When images should be sent, when IPs should start/end/reset
  - Establishing communication PC → FPGA-PS → FPGA-PL
- Ethernet was difficult to implement on the KC705/VC707
  - A lot of time was spent on this before switching to the ZCU104
- LWIP echo server requires a lot of networking knowledge
  - ZCU104 is capable of 1 Gbps transmission speeds
  - We hit a max of 500 Mbps transmission speeds
  - On average, our application has +/- 50 Mbps transmission speeds

# Challenges

- BRAM limitations constrained us
  - Currently working with  $\frac{1}{4}$  sized frame (448x232)
- Image preprocessing was tricky
  - Convert each pixel to IEEE 754 floating-point format
  - Started as a Python script, migrated to the FPGAs ARM core
- Floating-point functions took time
  - We created a library of floating-point arithmetic operations
- Each analysis mode in DICE has a lot of user-defined parameters and features

# Challenges

- Validation has been very difficult
  - Build source code from scratch with required libraries
  - Trace through dozens of functions and loops and manually cross-check numbers to results on FPGA
- How to accurately place a subset within the image
  - Current design requires manually listed pixel numbers for squares and circles
- The DICE source code is **extremely** dense
  - 98 total source code files that are being updated every week

# Overview

- Objective
- Results
- Challenges
- **Additional & Future Works**
- Conclusion
- Project Decomposition
  - Phase 1: Understanding DICE
  - Phase 2: Identifying Key Functions
  - Phase 3: Architecture Design
  - Phase 4: Client [PC]
  - Phase 5: Server [FPGA - PS]
  - Phase 6: Application [FPGA - PL]
  - Phase 7: Testing & Improvements

# Additional Work

- Additional work that is needed to make the project fully functional
  - Achieve at least 1 Gbps Ethernet transmission speeds
  - Add support for the missing results (beta, gamma, sigma)
  - Add support for the additional optimization/correlation method (simplex/robust)
  - Improved user-defined parameter support
  - Validation

# Future Work

- Addition of all analysis modes (subset-based full field, global)
- Image obstructions
- Full sized frame support (896x464)
- Faster Ethernet (10Gb+), or port to HMC
- Direct link between camera and FPGA (video → frames)
- Subset improvements (greater size, greater quantity, extra shapes, improved selection) [current design meets requirements]
- Add parallelization and general optimizations within code

# Overview

- Objective
- Results
- Challenges
- Additional & Future Works
- **Conclusion**
- Project Decomposition
  - Phase 1: Understanding DICE
  - Phase 2: Identifying Key Functions
  - Phase 3: Architecture Design
  - Phase 4: Client [PC]
  - Phase 5: Server [FPGA - PS]
  - Phase 6: Application [FPGA - PL]
  - Phase 7: Testing & Improvements

# Conclusion

- Implemented an end-to-end hardware accelerator system for DICE
  - Performed correlation on up to 1,000 frames with multiple subsets, 40,000 frames with a single subset
- 100% BRAM utilization (38 Mb total)
- Approximately 6,500 lines of Verilog written
- Complete design, code, and documentation is available on GitHub
  - [github.com/keatenstokke/DICE](https://github.com/keatenstokke/DICE)
- Development on this project will continue
  - Verify & Optimize



# Thank You Again, Honeywell

Questions?



**Honeywell**



# Overview

- Objective
- Results
- Challenges
- Additional & Future Works
- Conclusion
- **Project Decomposition**
  - **Phase 1: Understanding DICE**
  - Phase 2: Identifying Key Functions
  - Phase 3: Architecture Design
  - Phase 4: Client [PC]
  - Phase 5: Server [FPGA - PS]
  - Phase 6: Application [FPGA - PL]
  - Phase 7: Testing & Improvements

# Phase 1: Understanding DICE

- Traced through code to identify key functions and how they operate
  - Started with main.cpp and worked our way down
  - Hard to identify which functions belonged to the “tracking” analysis mode
- Code is **very dense**
  - Thousands of lines of code spread across 98 files
  - Lack of comments in the source code
  - No descriptions for functions or algorithms
- Building from the source code files was a hassle
  - Required the Trilinos library - contains 57 separate packages
  - Build instructions were not detailed, extra software required

# Overview

- Objective
- Results
- Challenges
- Additional & Future Works
- Conclusion
- **Project Decomposition**
  - Phase 1: Understanding DICE
  - **Phase 2: Identifying Key Functions**
  - Phase 3: Architecture Design
  - Phase 4: Client [PC]
  - Phase 5: Server [FPGA - PS]
  - Phase 6: Application [FPGA - PL]
  - Phase 7: Testing & Improvements

# Phase 2: Identifying Key Functions

## Main Functions

**main() → execute\_correlation() → generic\_correlation\_routine() →  
computeUpdateFast() → initial\_guess() → initialize\_guess\_4() →  
initialize() → interpolate\_bilinear() → interpolate\_grad\_x\_bilianer()  
& interpolate\_grad\_y\_bilianer() → gamma\_() → mean() →  
residuals\_aff() → map\_to\_u\_v\_theta\_aff() → map\_aff() →  
test\_for\_convergence\_aff() → save\_files()**

# Phase 2: Identifying Key Functions

## Main Functions

- **computeUpdateFast()**: Gradient-based optimization algorithm is done using a loop over the whole image to find the subset in the deformed image.
- **initial\_guess()**: Called to start the optimization with a good first guess.
- **initialize\_guess\_4()**: Initializes the subset in the vicinity of the previous guess.
- **initialize()**: Initializes the work variables.
- **interpolate\_bilinear()**: The easiest of ways to interpolate the pixel intensity values for non-pixel locations. Bilinear interpolation uses values of only the 4 nearest pixels, located in diagonal directions from a given pixel, in order to find the appropriate color intensity values of that pixel and tries to achieve a best approximation of a pixel's intensity based on the values at surrounding pixels.

# Phase 2: Identifying Key Functions

## Main functions

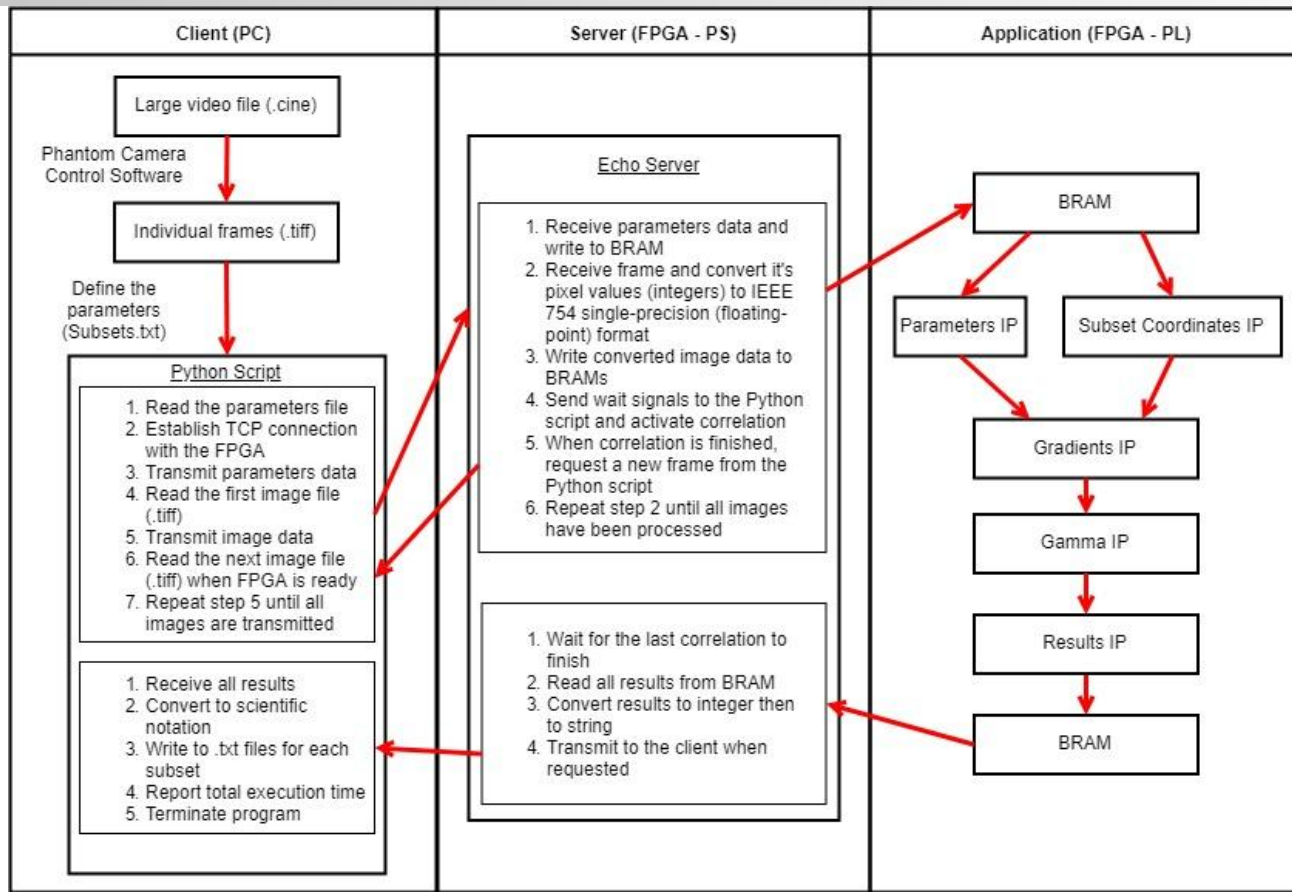
- **interpolate\_grad\_x\_bilinear() & interpolate\_grad\_y\_bilinear()**: The gradients also need to be interpolated when the values are requested at non-integer locations.
- **gamma\_()**: Returns the gamma correlation value between the reference and deformed subsets. It has a for-loop over the number of pixels per subset and computes the difference between the intensities of the pixels and the mean value of the intensities for that subset.
- **mean()**: Computes the mean intensity value for the reference or deformed subset.
- **residuals\_aff()**: Computes the residuals for the shape function (affine). Residuals are a measure of mismatch.
- **map\_to\_u\_v\_theta\_aff()**: Converts the current map parameters to u, v, and theta.
- **map\_aff()**: Maps the input coordinates to the output coordinates.
- **test\_for\_convergence\_aff()**: Returns true if the solution is converged.
- **save\_files()**: Saves off the parameters to the correct fields.

# Overview

- Objective
- Results
- Challenges
- Additional & Future Works
- Conclusion
- **Project Decomposition**
  - Phase 1: Understanding DICE
  - Phase 2: Identifying Key Functions
  - **Phase 3: Architecture Design**
  - Phase 4: Client [PC]
  - Phase 5: Server [FPGA - PS]
  - Phase 6: Application [FPGA - PL]
  - Phase 7: Testing & Improvements



# Phase 3: Architecture Design



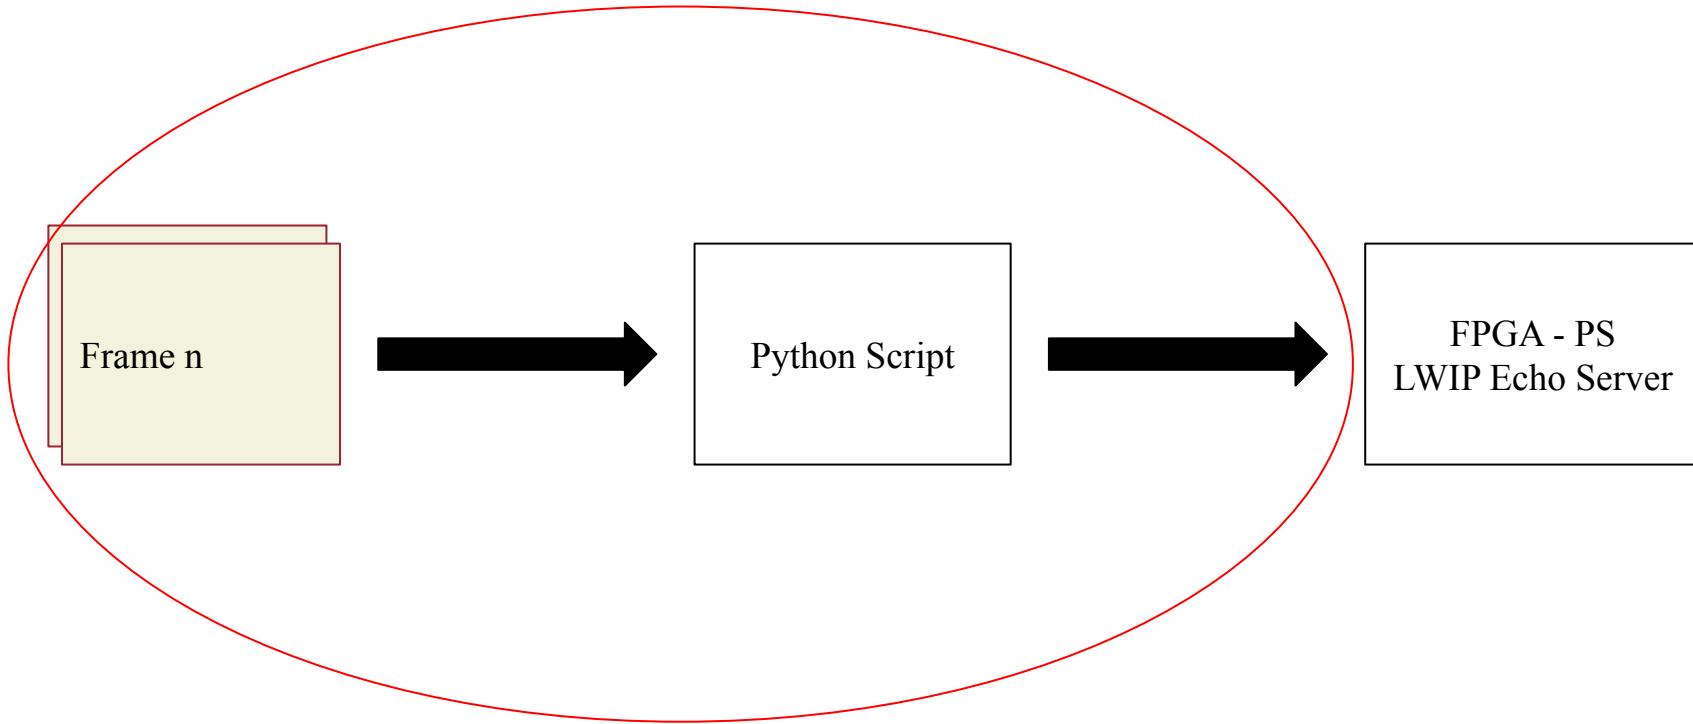
# Phase 3: Architecture Design

- Client used packets to send signals to communicate with the server
  - Server reciprocated with the proper response
  - End of frame, wait, send new frame, process done, results
- Server performed read and writes to memory within the PL
  - IPs wrote to slave registers to communicate with the server
  - AXI slave registers, BRAM controllers
- Application IPs communicated directly with each other via wires
  - Block diagrams used to connect IPs together to send data

# Overview

- Objective
- Results
- Challenges
- Additional & Future Works
- Conclusion
- **Project Decomposition**
  - Phase 1: Understanding DICE
  - Phase 2: Identifying Key Functions
  - Phase 3: Architecture Design
  - **Phase 4: Client [PC]**
  - Phase 5: Server [FPGA - PS]
  - Phase 6: Application [FPGA - PL]
  - Phase 7: Testing & Improvements

# Phase 4: Client [PC]



# Phase 4: Client [PC]

- A PC-based script was required to send the image data to the FPGA
- Python 2.7.15 was the language used
  - Great for prototyping - do more with less code
  - Versatile, easy to use, fast to develop
  - Lots of libraries available (Numpy, PIL, Bits, Codecs)
- The Python script acts as a TCP client
  - TCP is slower than UDP, but more reliable
  - We can't afford the loss of even 1 pixel

# Phase 4: Client [PC]

- 1) User runs script from terminal
  - a) All frames are in a specified directory
  - b) User-defined parameters are set (# of subsets, size, shape)
- 2) TCP connection is established with the FPGA
- 3) All parameters data is transmitted, begin looping through images
  - a) Packet size is 1,344 bytes, 3 bytes pixel, 448 pixel per packet
  - b) A total of 232 packets are sent for one 448x232 frame ( $\frac{1}{4}$  size)
- 4) Transmit each frame, wait until the server requests a new frame
- 5) Receive all results at the end of correlation, format and write to a file

# Phase 4: Client [PC]

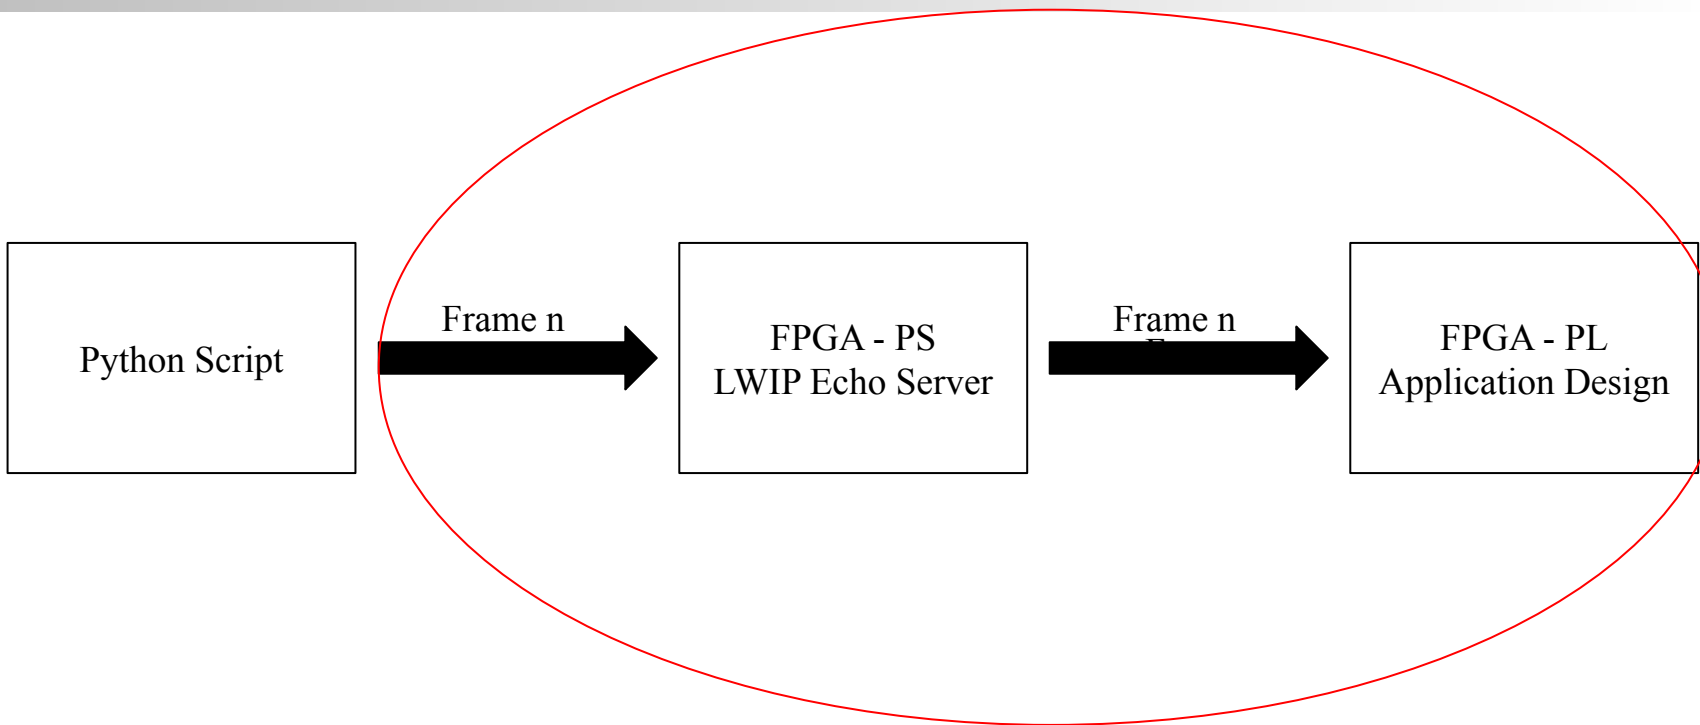
- Originally the Python script was going to read in the (.cine) video file and split into individual frames (.tif)
  - The Phantom cameras have free software that already does this
- The original script converted the pixels from integer to IEEE 754
  - We moved this process to the FPGA and saw a 5.33x speedup
  - Sending 4x less data to FPGA, less time in loops processing pixel values
  - FPGA dedicates an ARM core to processing this

# Overview

- Objective
- Results
- Challenges
- Additional & Future Works
- Conclusion
- **Project Decomposition**
  - Phase 1: Understanding DICE
  - Phase 2: Identifying Key Functions
  - Phase 3: Architecture Design
  - Phase 4: Client [PC]
  - **Phase 5: Server [FPGA - PS]**
  - Phase 6: Application [FPGA - PL]
  - Phase 7: Testing & Improvements



# Phase 5: Server [FPGA - PS]



# Phase 5: Server [FPGA - PS]

- Lightweight IP (LWIP) is a full TCP/IP networking stack for embedded systems
  - Other options: PetaLinux, FreeRTOS
- The echo server provided us with a quick template to receive and transmit data between the PC and the FPGA
  - We made modifications to the “recv\_callback” function to accept packets of predetermined length and transmit results
- The pixel data is received, converted to IEEE 754 floating-point format, then written into BRAM so the application can start

# Phase 5: Server [FPGA - PS]

- ZCU104 Ethernet transmission speeds
  - Capable of about 950 Mbps (1 Gb) speeds
  - Max transmission speed we were able to obtain was 500 Mbps (with buffer overflow errors)
  - Current design works with about +/- 50 Mbps speeds

Hardware Design Name	RAW Mode		Socket Mode	
	RX (Mb/s)	TX (Mb/s)	RX (Mb/s)	TX (Mb/s)
KC705_AxiEth_64kb_Cache	380	250	58.4	69.5
KC705_AxiEthernetlite_64kb_Cache	46	67	29	44
ZC702_GigE	943	949	521	542

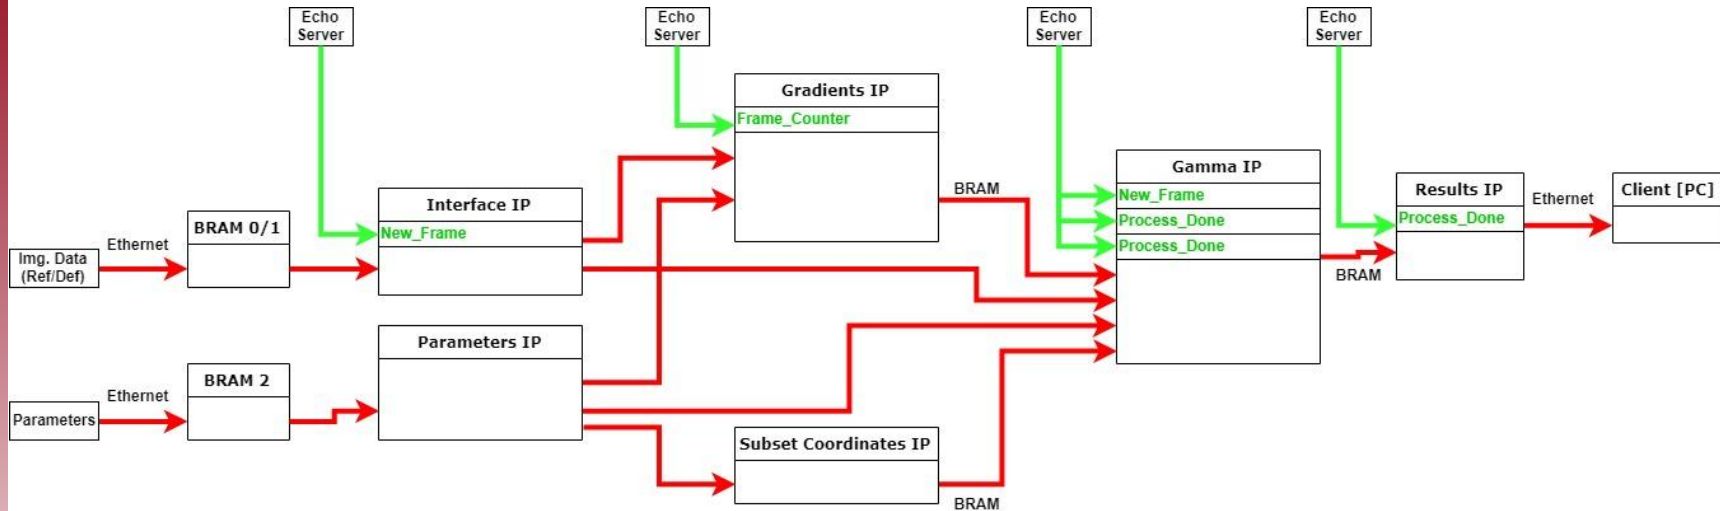
\* Table from  
Xilinx  
Reference  
Guide

# Overview

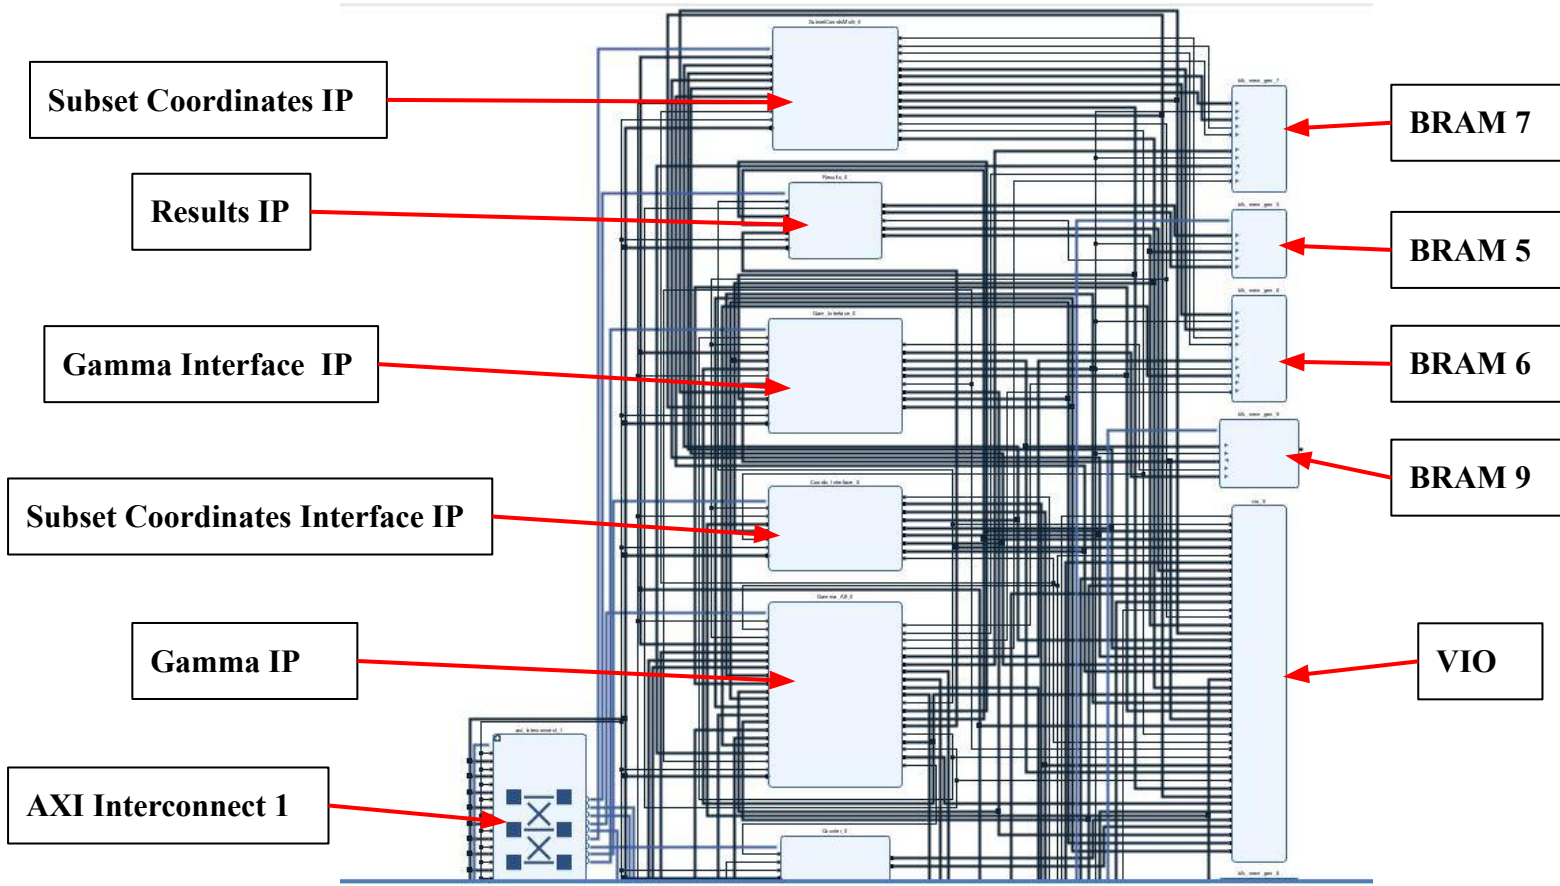
- Objective
- Results
- Challenges
- Additional & Future Works
- Conclusion
- **Project Decomposition**
  - Phase 1: Understanding DICE
  - Phase 2: Identifying Key Functions
  - Phase 3: Architecture Design
  - Phase 4: Client [PC]
  - Phase 5: Server [FPGA - PS]
  - **Phase 6: Application [FPGA - PL]**
  - Phase 7: Testing & Improvements

# Phase 6: Application [FPGA - PL]


**Simplified Block Diagram**



# Phase 6: Application [FPGA - PL]



# Phase 6: Application [FPGA - PL]



# Phase 6: Application [FPGA - PL]

- BRAMs/BRAM Controllers
  - Used for storing the local data (frames, gradients, subset coordinates, and user-defined parameters)
  - BRAM controllers interface with the echo server
- Parameters IP
  - Reads the user-defined parameters from BRAMs and sends them to the other IPs
  - Variables such as: number of subsets, shape, size, centerpoint, correlation method



# Phase 6: Application [FPGA - PL]

- Subset Coordinates IP
  - Computes the coordinates of the pixels that are inside the region of interest (subset) and saves the results into BRAMs
  - Tells the correlation where the subset is within a frame
- Subset Interface IP
  - Is responsible for reading the subset information from BRAMs and sending them to the Subset Coordinates IP as needed
  - Handles receiving and sending of data for multiple subsets

# Phase 6: Application [FPGA - PL]

- Interface IP
  - Handles the addressing to switch between the reference and the deformed frames in BRAM and sends the data to both the Gradients IP and Gamma IP
- Gradients IP
  - Gradients are the intensity between adjacent pixels
  - Computes the difference between intensities in both X and Y direction and saves the results into BRAMs

# Phase 6: Application [FPGA - PL]

- Gamma Interface IP
  - Is responsible for reading the subset information from BRAMs and sending them to the Gamma IP as needed
- Gamma IP
  - Uses the subset coordinates and gradients to perform the image correlation algorithm and sends the displacements to the Results IP
- Results IP
  - Receives the displacements from the Gamma IP and saves them into BRAMs
  - Interfaces with the echo server to send the results to the client

# Overview

- Objective
- Results
- Challenges
- Additional & Future Works
- Conclusion
- **Project Decomposition**
  - Phase 1: Understanding DICE
  - Phase 2: Identifying Key Functions
  - Phase 3: Architecture Design
  - Phase 4: Client [PC]
  - Phase 5: Server [FPGA - PS]
  - Phase 6: Application [FPGA - PL]
  - **Phase 7: Testing & Improvements**

# Phase 7: Testing

- Ethernet throughput tests
  - Sent 1 GB of data to the echo server to test transmission speeds
- Converting pixels to IEEE 754
  - Originally on Python script, added multiprocessing
  - Tested 1,000 numbers for correctness
  - Moved this step to the FPGA [ARM - PS]
- Performed unit testing on each function within the IPs
  - Used the VIO to monitor signals as they changed
  - Ran large scale tests with thousands of frames<sub>61</sub>

# Phase 7: Improvements

- Added support for multiple frames and multiple subsets
  - First tests were between 2 frames with 1 small subset
  - Added support for square and circular subsets
- Added support for user-defined parameters
  - Users can define the number of subsets, size, and shape prior to run
- Scaled up images from 64x48 to 448x232
  - Max size with our BRAM limitations;  $\frac{1}{4}$  full-frame
  - Moved from static images to real test images